

August 1, 2024

0.1 LAB ASSIGNMENT 02

Name: Imaad Hajwnae

Year: 4th

SRN: 202101132 / 23

Topic: Implementation of:

1. Multiple Regression Model
2. Naive Bayes - Classification & Regression
3. Random Forest Classification
4. Decision Tree Regression

Import Libraries

```
[1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns

from prettytable import PrettyTable
import missingno as msno
import matplotlib.pyplot as plt

data = pd.read_csv("C:
↪\\Users\\iamim\\OneDrive\\Desktop\\Seventh_Semester\\ML_LAB\\L2\\df_arabica_clean.
↪csv")
```

COLOR PALLET INITIALIZATION

```
[22]: purple = ['#491D8B', '#6929C4', '#8A3FFC', '#A56EFF',
               '#BE95FF', '#CA96EC', '#A163CF', '#29066B',
               '#7D3AC1', '#AF4BCE', '#DB4CB2', '#EB548C',
               '#EC96E0', '#A2128E', '#E8D9F3', '#641811']

sns.palplot(purple, size = 2)
```



```
[30]: data = data.copy()
data.head(n=10).style.background_gradient(cmap="Purples_r")
```

```
[30]: <pandas.io.formats.style.Styler at 0x1e3d3525b90>
```

```
[31]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 207 entries, 0 to 206
```

```
Data columns (total 41 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	207 non-null	int64
1	ID	207 non-null	int64
2	Country of Origin	207 non-null	object
3	Farm Name	205 non-null	object
4	Lot Number	206 non-null	object
5	Mill	204 non-null	object
6	ICO Number	75 non-null	object
7	Company	207 non-null	object
8	Altitude	206 non-null	object
9	Region	205 non-null	object
10	Producer	206 non-null	object
11	Number of Bags	207 non-null	int64
12	Bag Weight	207 non-null	object
13	In-Country Partner	207 non-null	object
14	Harvest Year	207 non-null	object
15	Grading Date	207 non-null	object
16	Owner	207 non-null	object
17	Variety	201 non-null	object
18	Status	207 non-null	object
19	Processing Method	202 non-null	object
20	Aroma	207 non-null	float64
21	Flavor	207 non-null	float64
22	Aftertaste	207 non-null	float64
23	Acidity	207 non-null	float64
24	Body	207 non-null	float64
25	Balance	207 non-null	float64
26	Uniformity	207 non-null	float64
27	Clean Cup	207 non-null	float64
28	Sweetness	207 non-null	float64
29	Overall	207 non-null	float64
30	Defects	207 non-null	float64
31	Total Cup Points	207 non-null	float64
32	Moisture Percentage	207 non-null	float64
33	Category One Defects	207 non-null	int64
34	Quakers	207 non-null	int64

```

35 Color                207 non-null    object
36 Category Two Defects 207 non-null    int64
37 Expiration           207 non-null    object
38 Certification Body    207 non-null    object
39 Certification Address 207 non-null    object
40 Certification Contact 207 non-null    object
dtypes: float64(13), int64(6), object(22)
memory usage: 66.4+ KB

```

```
[32]: data.describe().T.style.background_gradient(cmap = "magma")
```

```
[32]: <pandas.io.formats.style.Styler at 0x1e3d2512310>
```

```
[26]: print("Totally there are {} null values in the dataset".format(df.isnull().
      ↪sum().sum()))
```

Totally there are 153 null values in the dataset

```
[34]: table = PrettyTable()
      table.field_names = ['Feature', 'Data Type']

      for column in data.columns:
          column_dtype = str(data[column].dtype)
          table.add_row([column, column_dtype])

      print(table)
```

Feature	Data Type
Unnamed: 0	int64
ID	int64
Country of Origin	object
Farm Name	object
Lot Number	object
Mill	object
ICO Number	object
Company	object
Altitude	object
Region	object
Producer	object
Number of Bags	int64
Bag Weight	object
In-Country Partner	object
Harvest Year	object
Grading Date	object
Owner	object
Variety	object

Status	object
Processing Method	object
Aroma	float64
Flavor	float64
Aftertaste	float64
Acidity	float64
Body	float64
Balance	float64
Uniformity	float64
Clean Cup	float64
Sweetness	float64
Overall	float64
Defects	float64
Total Cup Points	float64
Moisture Percentage	float64
Category One Defects	int64
Quakers	int64
Color	object
Category Two Defects	int64
Expiration	object
Certification Body	object
Certification Address	object
Certification Contact	object

+-----+

Data Visualization

```
[35]: # Set a style for the plot
sns.set_style('whitegrid')

# Create a figure and axis object with a specific size
plt.figure(figsize=(14, 8))

# Generate the count plot with improved aesthetics
sns.countplot(
    data=df,
    y='Country of Origin',
    order=df['Country of Origin'].value_counts().index,
    palette='viridis' # Use a color palette for better aesthetics
)

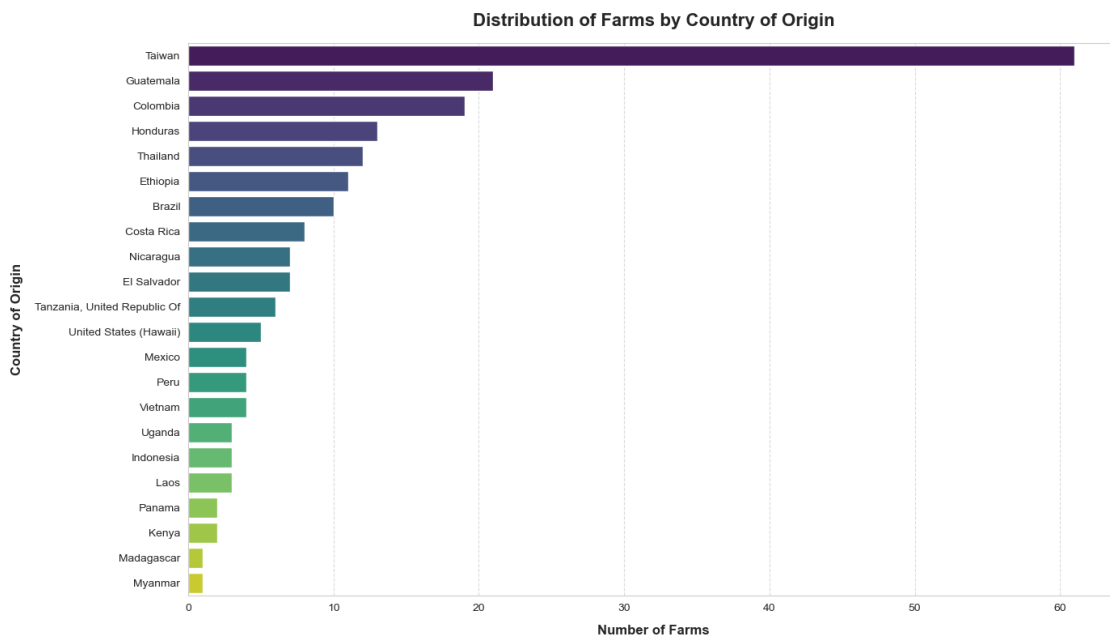
# Add labels and title with improved fonts and styles
plt.xlabel('Number of Farms', fontsize=12, fontweight='bold', labelpad=10)
plt.ylabel('Country of Origin', fontsize=12, fontweight='bold', labelpad=10)
plt.title('Distribution of Farms by Country of Origin', fontsize=16,
    fontweight='bold', pad=15)

# Add a grid for better readability
```

```
plt.grid(axis='x', linestyle='--', alpha=0.7)

# Adjust layout for better spacing
plt.tight_layout()

# Display the plot
plt.show()
```



```
[29]: # Calculate average points for each country
average_points = data.groupby('Country of Origin')['Total Cup Points'].mean().
    ↪reset_index()

# Sort countries by average total cup points in descending order
sorted_countries = average_points.sort_values('Total Cup Points',
    ↪ascending=False)

# Set the style and color palette for the plot
sns.set_style('whitegrid')
palette = sns.color_palette('Blues_d', len(sorted_countries))

# Create the figure and axis
plt.figure(figsize=(14, 10))

# Generate the bar plot
sns.barplot(
    data=sorted_countries,
```

```

    x='Total Cup Points',
    y='Country of Origin',
    palette=palette
)

# Set the labels and title with enhanced fonts
plt.xlabel('Average Total Cup Points', fontsize=12, fontweight='bold',
    ↪labelpad=10)
plt.ylabel('Country of Origin', fontsize=12, fontweight='bold', labelpad=10)
plt.title('Top-Rated Coffee Countries by Average Total Cup Points',
    ↪fontsize=16, fontweight='bold', pad=15)

# Add a grid for better readability
plt.grid(axis='x', linestyle='--', alpha=0.7)

# Customize ticks
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)

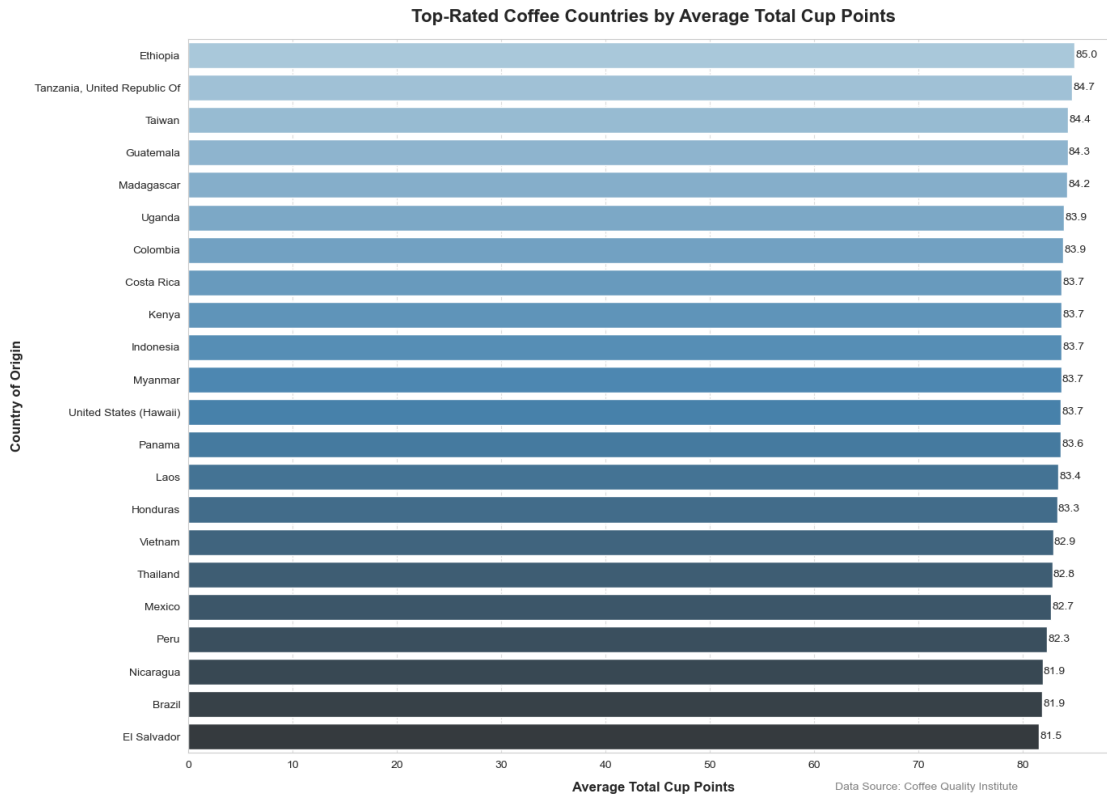
# Annotate bars with the average points value
for index, value in enumerate(sorted_countries['Total Cup Points']):
    plt.text(value + 0.1, index, f'{value:.1f}', va='center', fontsize=10)

# Add a legend or additional text if necessary
plt.figtext(0.9, 0.02, 'Data Source: Coffee Quality Institute',
    ↪horizontalalignment='right', fontsize=10, color='gray')

# Adjust layout to prevent clipping of ylabel/title
plt.tight_layout()

# Display the plot
plt.show()

```



```
[37]: import plotly.express as px

# Group by 'Country of Origin' and 'Farm Name', calculating the mean of 'Total Cup Points'
average_points = data.groupby(['Country of Origin', 'Farm Name'])['Total Cup Points'].mean().reset_index()

# Sort farms within each country by total cup points in descending order
sorted_farms = average_points.sort_values(['Country of Origin', 'Total Cup Points'], ascending=[True, False])

# Select the top farm for each country
top_farms = sorted_farms.groupby('Country of Origin').first().reset_index()

# Create a scatter plot using Plotly Express
fig = px.scatter(
    top_farms,
    x='Total Cup Points',
    y='Country of Origin',
    text='Farm Name',
    title='Top Coffee Farms in Each Country',
    color='Total Cup Points',
```

```

labels={'Total Cup Points': 'Average Total Cup Points'},
size='Total Cup Points',
size_max=15,
hover_data={
    'Total Cup Points': ':.2f', # Format Total Cup Points to two decimal
    ↪places
    'Country of Origin': True,
    'Farm Name': True
},
template='plotly_white'
)

# Update traces for better text position and marker aesthetics
fig.update_traces(
    textposition='top center',
    marker=dict(
        line=dict(width=1, color='DarkSlateGrey'),
        opacity=0.7
    )
)

# Update layout with enhanced title, axis labels, and grid
fig.update_layout(
    xaxis_title='Average Total Cup Points',
    yaxis_title='Country of Origin',
    title={
        'text': 'Top Coffee Farms in Each Country by Average Total Cup Points',
        'font': {'size': 20, 'family': 'Arial', 'color': 'DarkSlateGrey'},
        'x': 0.5, # Center the title
        'xanchor': 'center'
    },
    xaxis=dict(
        tickfont=dict(size=12),
        titlefont=dict(size=14, family='Arial', color='Black'),
        gridcolor='LightGray',
        zerolinecolor='Gray'
    ),
    yaxis=dict(
        tickfont=dict(size=12),
        titlefont=dict(size=14, family='Arial', color='Black'),
        gridcolor='LightGray',
        zerolinecolor='Gray'
    ),
    coloraxis_colorbar=dict(
        title='Total Cup Points',
        thickness=15,
        len=0.5,

```



```

        yanchor='middle',
        xanchor='right',
        xpad=20
    ),
    hoverlabel=dict(
        bgcolor='white',
        font_size=12,
        font_family='Arial'
    ),
    margin=dict(l=60, r=60, t=80, b=60), # Adjust margins
    width=900,
    height=600
)

# Show the plot
fig.show()

```

```

[38]: columns = ['Flavor', 'Aroma', 'Aftertaste', 'Acidity', 'Body', 'Balance',
                'Uniformity', 'Sweetness', 'Overall']
data['Mean Score'] = data[columns].mean(axis=1)
best_coffee_row = data.loc[data['Mean Score'].idxmax()]
best_coffee_details = best_coffee_row[['Company', 'Variety', 'Country of_
    ↪Origin', 'Farm Name', 'Processing Method']]
print("Best Coffee Combination:")
print("Company:", best_coffee_details['Company'])
print("Variety:", best_coffee_details['Variety'])
print("Country of Origin:", best_coffee_details['Country of Origin'])
print("Farm Name:", best_coffee_details['Farm Name'])
print("Processing Method:", best_coffee_details['Processing Method'])

```

```

Best Coffee Combination:
Company: Coffee Quality Union
Variety: Castillo
Country of Origin: Colombia
Farm Name: Finca El Paraiso
Processing Method: Double Anaerobic Washed

```

```

[39]: import pandas as pd
import plotly.express as px

# Step 1: Calculate Farm Count per Company
company_counts = data['Company'].value_counts().reset_index()
company_counts.columns = ['Company', 'Farm Count']

# Step 2: Sort Companies by Farm Count in Descending Order
sorted_companies = company_counts.sort_values('Farm Count', ascending=False)

```

```

# Step 3: Select Top 5 Companies with the Most Farms
top_companies = sorted_companies.head(5)

# Step 4: Analyze Moisture Percentage for Each Variety of Coffee in Top
↳ Companies
for company in top_companies['Company']:
    company_data = data[data['Company'] == company]

    # Calculate the average moisture percentage per variety for the company
    variety_moisture = company_data.groupby('Variety')['Moisture Percentage'].
    ↳ mean().reset_index()

    # Sort varieties by moisture percentage in ascending order
    sorted_varieties = variety_moisture.sort_values('Moisture Percentage',
    ↳ ascending=True)

    # Identify the variety with the best (lowest) moisture percentage
    best_variety = sorted_varieties.iloc[0]['Variety']

# Step 5: Plot the Moisture Percentage for Each Variety
fig = px.bar(
    sorted_varieties,
    x='Variety',
    y='Moisture Percentage',
    title=f'Average Moisture Percentage by Variety for {company}',
    labels={'Moisture Percentage': 'Average Moisture Percentage'},
    text='Moisture Percentage',
    template='plotly_white',
    color='Moisture Percentage',
    color_continuous_scale='Blues'
)

# Enhance the figure with additional layout and styling
fig.update_traces(
    texttemplate='%{text:.2f}%', # Format the moisture percentage to two
    ↳ decimal places
    textposition='outside',      # Place the text labels outside the bars
    marker=dict(
        line=dict(width=1, color='DarkSlateGrey')
    )
)

fig.update_layout(
    xaxis_title='Variety',
    yaxis_title='Average Moisture Percentage',
    title={
        'text': f'Average Moisture Percentage by Variety for {company}',

```

```

        'font': {'size': 20, 'family': 'Arial', 'color': 'DarkSlateGrey'},
        'x': 0.5, # Center the title
        'xanchor': 'center'
    },
    xaxis=dict(
        tickfont=dict(size=12),
        titlefont=dict(size=14, family='Arial', color='Black'),
        gridcolor='LightGray',
        zerolinecolor='Gray'
    ),
    yaxis=dict(
        tickfont=dict(size=12),
        titlefont=dict(size=14, family='Arial', color='Black'),
        gridcolor='LightGray',
        zerolinecolor='Gray'
    ),
    margin=dict(l=60, r=60, t=80, b=60), # Adjust margins
    width=900,
    height=600
)

# Display the plot
fig.show()

```

[]:

```

[40]: company_counts = data['Company'].value_counts().reset_index()
company_counts.columns = ['Company', 'Farm Count']
sorted_companies = company_counts.sort_values('Farm Count', ascending=False)
top_companies = sorted_companies.head(5)
for company in top_companies['Company']:
    company_data = data[data['Company'] == company]
    variety_flavor = company_data.groupby('Variety')['Flavor'].mean().
    ↪reset_index()
    sorted_varieties = variety_flavor.sort_values('Flavor', ascending=False)
    best_variety = sorted_varieties.iloc[0]['Variety']
    fig = px.bar(sorted_varieties, x='Variety', y='Flavor',
                  title=f'Best Flavor Variety for {company}', labels={'Flavor': '
    ↪Average Flavor'})

    fig.show()

```

```

[41]: variety_counts = data['Variety'].value_counts().reset_index()
variety_counts.columns = ['Variety', 'Farm Count']
sorted_varieties = variety_counts.sort_values('Farm Count', ascending=False)
top_varieties = sorted_varieties.head(15)
fig = px.bar(top_varieties, x='Variety', y='Farm Count',

```

```

        title='Top Varieties', labels={'Farm Count': 'Number of Farms'})
fig.show()
fig = px.sunburst(top_varieties, path=['Variety'], values='Farm Count',
                  title='Top Varieties')
fig.show()

```

```

[42]: import pandas as pd
import plotly.express as px

# Step 1: Calculate the Number of Farms per Company
company_counts = data['Company'].value_counts().reset_index()
company_counts.columns = ['Company', 'Farm Count']

# Step 2: Sort Companies by Farm Count in Descending Order
sorted_companies = company_counts.sort_values('Farm Count', ascending=False)

# Step 3: Identify the Top Company with the Most Farms
top_company = sorted_companies.iloc[0]['Company']

# Step 4: Filter Data for the Top Company
top_company_data = data[data['Company'] == top_company]

# Step 5: Calculate the Number of Farms per Country for the Top Company
country_counts = top_company_data['Country of Origin'].value_counts().
    ↪reset_index()
country_counts.columns = ['Country', 'Farm Count']

# Step 6: Create a Choropleth Map to Show Farm Distribution by Country
fig = px.choropleth(
    country_counts,
    locations='Country',
    locationmode='country names',
    color='Farm Count',
    title=f'Distribution of Farms by Country for Top Company: {top_company}',
    labels={'Farm Count': 'Number of Farms'},
    color_continuous_scale=px.colors.sequential.Plasma,
    template='plotly_white',
    hover_data={'Country': True, 'Farm Count': True}
)

# Step 7: Enhance the Layout and Aesthetics of the Map
fig.update_geos(
    projection_type='natural earth', # Use a natural earth projection
    showcountries=True,
    countrycolor='LightGrey',
    showcoastlines=True,
    coastlinecolor='LightBlue',

```

```

        showland=True,
        landcolor='whitesmoke'
    )

fig.update_layout(
    title={
        'text': f'Distribution of Farms by Country for Top Company:␣
↪{top_company}',
        'font': {'size': 20, 'family': 'Arial', 'color': 'DarkSlateGrey'},
        'x': 0.5, # Center the title
        'xanchor': 'center'
    },
    margin=dict(l=50, r=50, t=100, b=50), # Adjust margins
    coloraxis_colorbar=dict(
        title='Number of Farms',
        thickness=15,
        len=0.5,
        yanchor='middle',
        xanchor='right',
        xpad=20
    ),
    hoverlabel=dict(
        bgcolor='white',
        font_size=12,
        font_family='Arial'
    ),
    width=1000,
    height=600
)

# Step 8: Display the Choropleth Map
fig.show()

```

```

[43]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Define the list of variables for analysis
variables = [
    'Aroma', 'Flavor', 'Aftertaste', 'Acidity', 'Body', 'Balance',
    'Uniformity', 'Clean Cup', 'Sweetness', 'Overall', 'Defects',
    'Total Cup Points', 'Moisture Percentage', 'Category One Defects',
    'Quakers'
]

# Set the style for the plots
sns.set(style='whitegrid', palette='muted', font_scale=1.2)

```

```

# Step 1: Histograms for Each Variable
plt.figure(figsize=(15, 12))
data[variables].hist(bins=10, figsize=(15, 12), color='skyblue',
    ↪edgecolor='black', grid=False)
plt.suptitle('Distribution of Coffee Quality Attributes', fontsize=20,
    ↪fontweight='bold')
plt.tight_layout(rect=[0, 0.03, 1, 0.95]) # Adjust layout to make space for
    ↪the title
plt.show()

# Step 2: Boxplot for Each Variable
plt.figure(figsize=(15, 10))
sns.boxplot(data=data[variables], palette='pastel')
plt.title('Boxplots of Coffee Quality Attributes', fontsize=20,
    ↪fontweight='bold')
plt.xlabel('Attributes', fontsize=15, fontweight='bold')
plt.ylabel('Values', fontsize=15, fontweight='bold')
plt.xticks(rotation=45, fontsize=12)
plt.yticks(fontsize=12)
plt.show()

# Step 3: Scatter Plot - Flavor vs. Total Cup Points
plt.figure(figsize=(12, 8))
sns.scatterplot(data=data, x='Flavor', y='Total Cup Points', hue='Country of
    ↪Origin',
    palette='coolwarm', s=100, alpha=0.8, edgecolor='w',
    ↪linewidth=0.8)
plt.title('Relationship Between Flavor and Total Cup Points', fontsize=20,
    ↪fontweight='bold')
plt.xlabel('Flavor', fontsize=15, fontweight='bold')
plt.ylabel('Total Cup Points', fontsize=15, fontweight='bold')
plt.legend(title='Country of Origin', bbox_to_anchor=(1.05, 1), loc='upper
    ↪left', fontsize=12)
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()

# Step 4: Bar Plot - Country of Origin vs. Total Cup Points
plt.figure(figsize=(15, 8))
sns.barplot(data=data, x='Country of Origin', y='Total Cup Points',
    ↪palette='viridis', ci=None)
plt.title('Average Total Cup Points by Country of Origin', fontsize=20,
    ↪fontweight='bold')
plt.xlabel('Country of Origin', fontsize=15, fontweight='bold')
plt.ylabel('Average Total Cup Points', fontsize=15, fontweight='bold')
plt.xticks(rotation=45, fontsize=12)

```

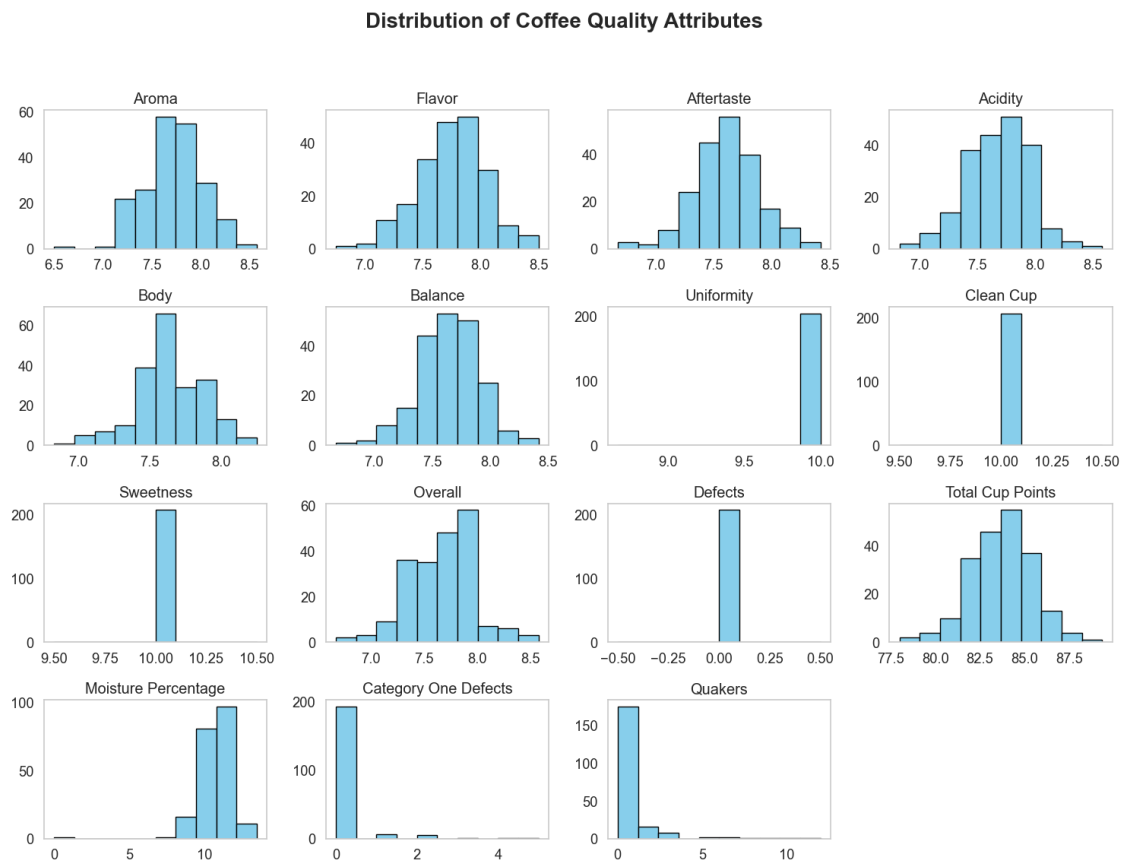
```

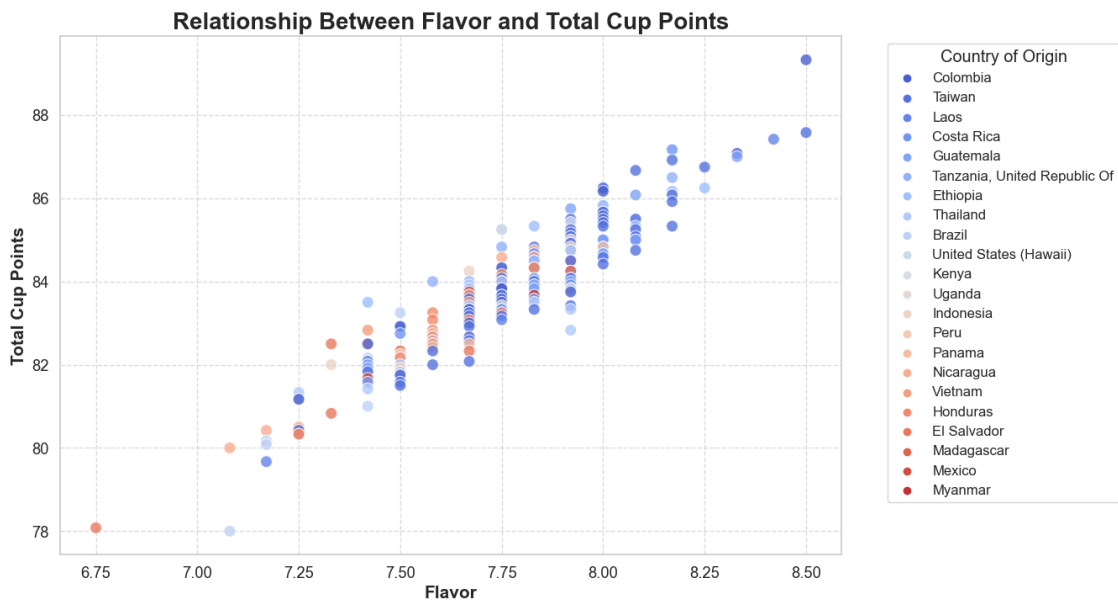
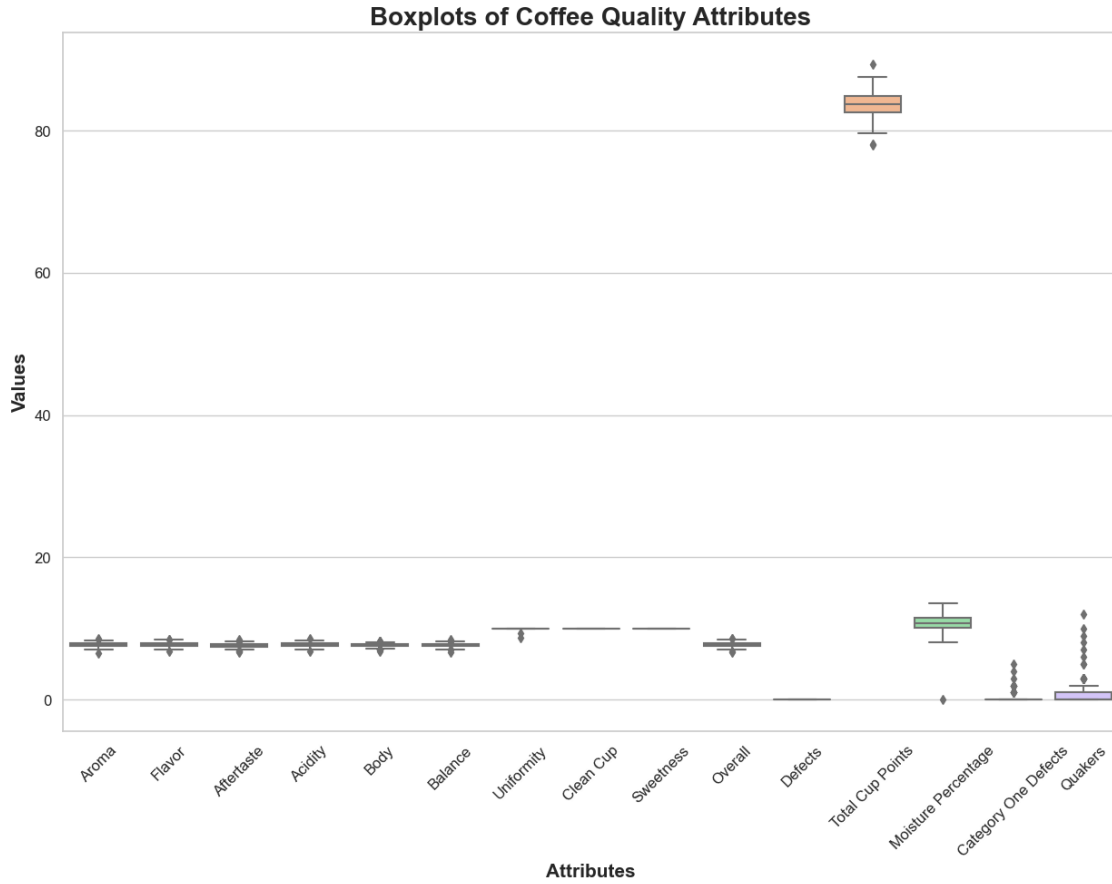
plt.yticks(fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()

# Step 5: Correlation Heatmap
plt.figure(figsize=(15, 10))
corr = data[variables].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5,
            annot_kws={'size': 10, 'weight': 'bold', 'color': 'darkblue'})
plt.title('Correlation Heatmap of Coffee Quality Attributes', fontsize=20,
          fontweight='bold')
plt.xticks(rotation=45, fontsize=12)
plt.yticks(fontsize=12)
plt.show()

```


<Figure size 1500x1200 with 0 Axes>

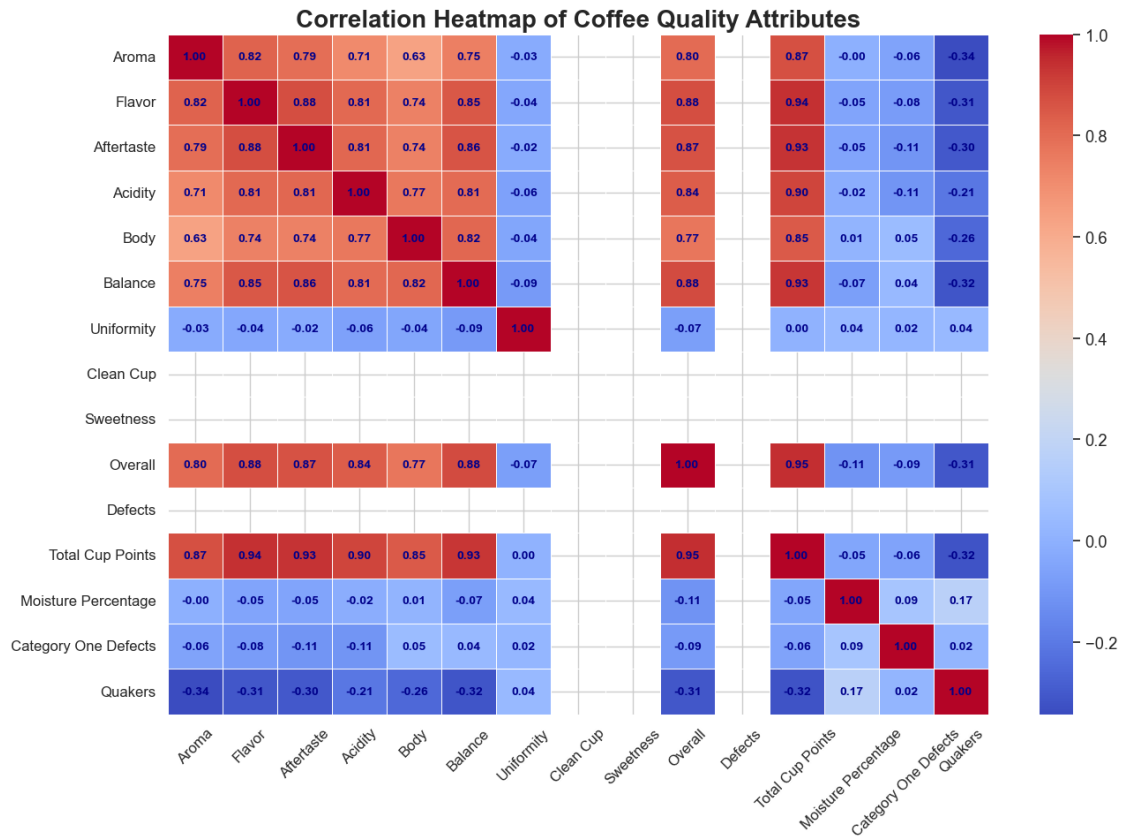




C:\Users\iamim\AppData\Local\Temp\ipykernel_17008\3081503852.py:46:
FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.





Correlation Analysis

```
[44]: import plotly.express as px

# Define the variables to analyze correlations
variables = [
    'Aroma', 'Flavor', 'Aftertaste', 'Acidity', 'Body', 'Balance',
    'Uniformity', 'Sweetness', 'Overall', 'Defects', 'Total Cup Points'
]

# Calculate the correlation matrix
corr_matrix = data[variables].corr()

# Create the heatmap using Plotly
fig = px.imshow(
    corr_matrix,
    labels=dict(x="Variables", y="Variables", color="Correlation"),
    x=variables,
    y=variables,
    color_continuous_scale='RdBu_r', # Reversed 'RdBu' for a more intuitive
    ↪ color mapping
```

```

    zmin=-1, # Set minimum correlation value for consistent scaling
    zmax=1,  # Set maximum correlation value for consistent scaling
    title='Correlation Heatmap of Coffee Quality Attributes',
    aspect="auto", # Ensures aspect ratio is automatically determined
)

# Update layout to enhance appearance
fig.update_layout(
    title={
        'text': 'Correlation Heatmap of Coffee Quality Attributes',
        'font': {'size': 24, 'family': 'Arial', 'color': 'DarkSlateGray'},
        'x': 0.5,
        'xanchor': 'center'
    },
    xaxis_title='Coffee Quality Attributes',
    yaxis_title='Coffee Quality Attributes',
    xaxis=dict(
        tickangle=-45,
        tickfont=dict(size=12, family='Arial', color='Black')
    ),
    yaxis=dict(
        tickfont=dict(size=12, family='Arial', color='Black')
    ),
    width=800,
    height=800,
    margin=dict(l=100, r=100, t=100, b=100),
    coloraxis_colorbar=dict(
        title='Correlation',
        thickness=15,
        len=0.75,
        ticks='outside',
        tickfont=dict(size=12),
        titlefont=dict(size=14)
    ),
    plot_bgcolor='white', # Background color of the plot
    hoverlabel=dict(
        bgcolor="white",
        font_size=12,
        font_family="Arial"
    ),
)

# Add annotations for better clarity
for i in range(len(corr_matrix)):
    for j in range(len(corr_matrix)):
        fig.add_annotation(
            x=variables[j],

```

```

        y=variables[i],
        text=f"{corr_matrix.iloc[i, j]:.2f}",
        showarrow=False,
        font=dict(size=11, color='black' if abs(corr_matrix.iloc[i, j]) < 0.
↪6 else 'white') # Adjust text color based on background
    )

# Show the figure
fig.show()

```

Identification of Numerical Variables in the dataset

```

[45]: import pandas as pd

# Assuming 'data' is your DataFrame
numerical_columns = data.select_dtypes(include=['number']).columns

# Print the names of numerical columns
print("Numerical Columns in the DataFrame:")
for col in numerical_columns:
    print(col)

```

Numerical Columns in the DataFrame:

```

Unnamed: 0
ID
Number of Bags
Aroma
Flavor
Aftertaste
Acidity
Body
Balance
Uniformity
Clean Cup
Sweetness
Overall
Defects
Total Cup Points
Moisture Percentage
Category One Defects
Quakers
Category Two Defects
Mean Score

```

0.1.1 Multiple Regression Variable Modeling

```
[47]: # Import required libraries
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Display the first few rows of the dataset
data.head()
```

```
[47]: Unnamed: 0  ID Country of Origin          Farm Name \
0           0    0          Colombia          Finca El Paraiso
1           1    1          Taiwan  Royal Bean Geisha Estate
2           2    2          Laos      OKLAO coffee farms
3           3    3      Costa Rica          La Cumbre
4           4    4          Colombia          Finca Santuario

          Lot Number          Mill \
0          CQU2022015          Finca El Paraiso
1  The 2022 Pacific Rim Coffee Summit,T037      Royal Bean Geisha Estate
2  The 2022 Pacific Rim Coffee Summit,LA01  oklao coffee processing plant
3          CQU2022017      La Montana Tarrazu Mill
4          CQU2023002          Finca Santuario

ICO Number          Company  Altitude          Region ... \
0          NaN      Coffee Quality Union  1700-1930      Piendamó,Cauca ...
1          NaN  Taiwan Coffee Laboratory      1200          Chiayi ...
2          NaN  Taiwan Coffee Laboratory      1300  Laos Borofen Plateau ...
3          NaN      Coffee Quality Union      1900  Los Santos,Tarrazu ...
4          NaN      Coffee Quality Union  1850-2100      Popayan,Cauca ...

Moisture Percentage  Category One Defects  Quakers          Color \
0           11.8           0           0          green
1           10.5           0           0      blue-green
2           10.4           0           0      yellowish
3           11.8           0           0          green
4           11.6           0           2  yellow-green

Category Two Defects          Expiration \
0           3  September 21st, 2023
1           0  November 15th, 2023
2           2  November 15th, 2023
3           0  September 21st, 2023
4           2    March 5th, 2024

          Certification Body \
0      Japan Coffee Exchange
```

```

1 Taiwan Coffee Laboratory
2 Taiwan Coffee Laboratory
3         Japan Coffee Exchange
4         Japan Coffee Exchange

                                Certification Address \
0  413-0002          - 1173-58 Izusan, Ata...
1  QAHWAH CO., LTD 4F, No. 225, Sec. 3, Beixin Rd...
2  QAHWAH CO., LTD 4F, No. 225, Sec. 3, Beixin Rd...
3  413-0002          - 1173-58 Izusan, Ata...
4  413-0002          - 1173-58 Izusan, Ata...

                                Certification Contact Mean Score
0      Koju Matsuzawa - +81(0)9085642901    8.814444
1      Lin, Jen-An Neil - 886-289116612    8.621111
2      Lin, Jen-An Neil - 886-289116612    8.602222
3      Koju Matsuzawa - +81(0)9085642901    8.574444
4      Koju Matsuzawa - +81(0)9085642901    8.564444

[5 rows x 42 columns]

```

```

[52]: # Display data types
print("Data Types:")
print(data.dtypes)

# Separate numeric and non-numeric columns
numeric_cols = data.select_dtypes(include=['number']).columns
non_numeric_cols = data.select_dtypes(exclude=['number']).columns

print("\nNumeric Columns:")
print(numeric_cols)

print("\nNon-Numeric Columns:")
print(non_numeric_cols)

# Fill missing values for numeric columns with column mean
data[numeric_cols] = data[numeric_cols].fillna(data[numeric_cols].mean())

# Handle missing values in non-numeric columns
for col in non_numeric_cols:
    data[col].fillna(data[col].mode()[0], inplace=True) # Example: filling
    ↳ with mode

# Verify missing values are handled
print("\nMissing Values after Imputation:")
print(data.isnull().sum())

```

```
# Verify data types after handling missing values
print("\nData Types after Imputation:")
print(data.dtypes)
```

Data Types:

Unnamed: 0	int64
ID	int64
Country of Origin	object
Farm Name	object
Lot Number	object
Mill	object
ICO Number	object
Company	object
Altitude	object
Region	object
Producer	object
Number of Bags	int64
Bag Weight	object
In-Country Partner	object
Harvest Year	object
Grading Date	object
Owner	object
Variety	object
Status	object
Processing Method	object
Aroma	float64
Flavor	float64
Aftertaste	float64
Acidity	float64
Body	float64
Balance	float64
Uniformity	float64
Clean Cup	float64
Sweetness	float64
Overall	float64
Defects	float64
Total Cup Points	float64
Moisture Percentage	float64
Category One Defects	int64
Quakers	int64
Color	object
Category Two Defects	int64
Expiration	object
Certification Body	object
Certification Address	object
Certification Contact	object
Mean Score	float64

dtype: object

Numeric Columns:

```
Index(['Unnamed: 0', 'ID', 'Number of Bags', 'Aroma', 'Flavor', 'Aftertaste',  
      'Acidity', 'Body', 'Balance', 'Uniformity', 'Clean Cup', 'Sweetness',  
      'Overall', 'Defects', 'Total Cup Points', 'Moisture Percentage',  
      'Category One Defects', 'Quakers', 'Category Two Defects',  
      'Mean Score'],  
      dtype='object')
```

Non-Numeric Columns:

```
Index(['Country of Origin', 'Farm Name', 'Lot Number', 'Mill', 'ICO Number',  
      'Company', 'Altitude', 'Region', 'Producer', 'Bag Weight',  
      'In-Country Partner', 'Harvest Year', 'Grading Date', 'Owner',  
      'Variety', 'Status', 'Processing Method', 'Color', 'Expiration',  
      'Certification Body', 'Certification Address', 'Certification Contact'],  
      dtype='object')
```

Missing Values after Imputation:

Unnamed: 0	0
ID	0
Country of Origin	0
Farm Name	0
Lot Number	0
Mill	0
ICO Number	0
Company	0
Altitude	0
Region	0
Producer	0
Number of Bags	0
Bag Weight	0
In-Country Partner	0
Harvest Year	0
Grading Date	0
Owner	0
Variety	0
Status	0
Processing Method	0
Aroma	0
Flavor	0
Aftertaste	0
Acidity	0
Body	0
Balance	0
Uniformity	0
Clean Cup	0
Sweetness	0

Overall	0
Defects	0
Total Cup Points	0
Moisture Percentage	0
Category One Defects	0
Quakers	0
Color	0
Category Two Defects	0
Expiration	0
Certification Body	0
Certification Address	0
Certification Contact	0
Mean Score	0

dtype: int64

Data Types after Imputation:

Unnamed: 0	int64
ID	int64
Country of Origin	object
Farm Name	object
Lot Number	object
Mill	object
ICO Number	object
Company	object
Altitude	object
Region	object
Producer	object
Number of Bags	int64
Bag Weight	object
In-Country Partner	object
Harvest Year	object
Grading Date	object
Owner	object
Variety	object
Status	object
Processing Method	object
Aroma	float64
Flavor	float64
Aftertaste	float64
Acidity	float64
Body	float64
Balance	float64
Uniformity	float64
Clean Cup	float64
Sweetness	float64
Overall	float64
Defects	float64
Total Cup Points	float64

```

Moisture Percentage    float64
Category One Defects   int64
Quakers               int64
Color                 object
Category Two Defects   int64
Expiration             object
Certification Body     object
Certification Address  object
Certification Contact  object
Mean Score            float64
dtype: object

```

```

[53]: # Now continue with feature selection, splitting, training, and evaluating the
      ↪ model
      independent_vars = ['Aroma', 'Flavor'] # Example independent variables
      dependent_var = 'Total Cup Points' # Example dependent variable

      # Extract features and target
      X = data[independent_vars]
      y = data[dependent_var]

```

```

[54]: # Split the data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)

```

```

[55]: # Initialize and train the Linear Regression model
      regression_model = LinearRegression()
      regression_model.fit(X_train, y_train)

```

```

[55]: LinearRegression()

```

```

[56]: # Display model coefficients
      print("\nModel Coefficients:")
      print(f"Intercept: {regression_model.intercept_}")
      for idx, coef in enumerate(regression_model.coef_):
          print(f"Coefficient for {independent_vars[idx]}: {coef}")

```

```

Model Coefficients:
Intercept: 35.79207032807622
Coefficient for Aroma: 1.9013565731881188
Coefficient for Flavor: 4.292202508997079

```

```

[57]: # Predict on the test set and evaluate the model
      y_pred = regression_model.predict(X_test)
      mse = mean_squared_error(y_test, y_pred)
      r2 = r2_score(y_test, y_pred)

```

```

print(f"\nModel Evaluation Metrics:")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R2 Score: {r2:.2f}")

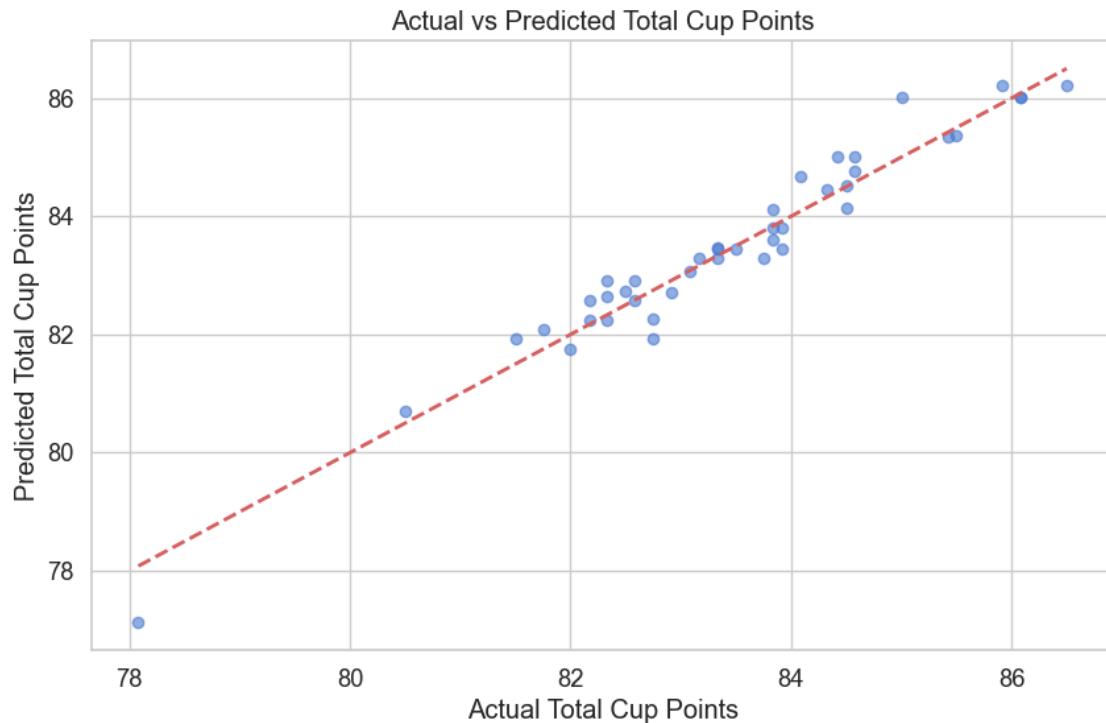
```

Model Evaluation Metrics:
 Mean Squared Error (MSE): 0.14
 R² Score: 0.94

```

[58]: # Plot actual vs predicted values
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.6, color='b')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel('Actual Total Cup Points')
plt.ylabel('Predicted Total Cup Points')
plt.title('Actual vs Predicted Total Cup Points')
plt.show()

```



0.1.2 Navie bayis

Classification

0.1.3

```
[63]: from sklearn.model_selection import train_test_split
      from sklearn.naive_bayes import GaussianNB
      from sklearn.metrics import accuracy_score, classification_report, \
          ↪confusion_matrix

      # Define features and target
      features = ['Aroma', 'Flavor', 'Aftertaste', 'Acidity', 'Body', 'Balance', \
          ↪'Uniformity',
                  'Clean Cup', 'Sweetness', 'Overall', 'Defects', 'Total Cup Points',
                  'Moisture Percentage', 'Category One Defects', 'Quakers', \
          ↪'Category Two Defects', 'Mean Score']
      target = 'Category One Defects' # Assuming this is categorical

      # Extract features and target
      X = data[features]
      y = data[target]

      # Check the unique values in the target variable
      print("\nUnique values in target variable:")
      print(y.unique())
```

Unique values in target variable:
[0 2 1 3 4 5]

```
[64]: # Split the data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
          ↪random_state=42)

      # Check the shapes of the train and test sets
      print(f"Training set shape (X_train): {X_train.shape}")
      print(f"Test set shape (X_test): {X_test.shape}")
```

Training set shape (X_train): (165, 17)
Test set shape (X_test): (42, 17)

```
[65]: # Initialize the Naive Bayes classifier
      nb_model = GaussianNB()

      # Train the model using the training data
      nb_model.fit(X_train, y_train)

      # Display model parameters (optional)
      print("\nModel Parameters:")
      print(nb_model.get_params())
```

Model Parameters:
{'priors': None, 'var_smoothing': 1e-09}

```
[66]: # Predict on the test set
y_pred = nb_model.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"\nAccuracy: {accuracy:.2f}")

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(conf_matrix)

# Plot confusion matrix using seaborn
plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=nb_model.classes_, yticklabels=nb_model.classes_)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```

Accuracy: 0.88

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	36
1	0.00	0.00	0.00	0
2	0.25	0.33	0.29	3
3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	1
5	0.00	0.00	0.00	1
accuracy			0.88	42
macro avg	0.21	0.22	0.21	42
weighted avg	0.88	0.88	0.88	42

Confusion Matrix:

```
[[36  0  0  0  0  0]
 [ 0  0  0  0  0  0]
 [ 0  2  1  0  0  0]
 [ 0  0  1  0  0  0]
 [ 0  0  1  0  0  0]
 [ 0  0  1  0  0  0]]
```

```
c:\Users\iamim\anaconda3\Lib\site-
```

```
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
```

```
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
```

```
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
```

```
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
```

```
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
```

```
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.



Regression

```
[73]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import BayesianRidge
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[74]: # Define features and target
features = ['Aroma', 'Flavor', 'Aftertaste', 'Acidity', 'Body', 'Balance',
           ↪ 'Uniformity',
           'Clean Cup', 'Sweetness', 'Overall', 'Defects', 'Total Cup Points',
           'Moisture Percentage', 'Category One Defects', 'Quakers',
           ↪ 'Category Two Defects', 'Mean Score']
target = 'Total Cup Points' # Change to your target variable for regression
X = data[features]
```

```
y = data[target]
```

```
[75]: # Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[76]: # Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
[77]: # Initialize and train Bayesian Ridge Regression model
model = BayesianRidge()
model.fit(X_train_scaled, y_train)
```

```
[77]: BayesianRidge()
```

```
[78]: # Predict on the test set
y_pred = model.predict(X_test_scaled)


# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R^2 Score: {r2:.2f}")

# Plot predictions vs actual values
plt.figure(figsize=(10, 6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.6)
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs Predicted Values')
plt.show()
```

Mean Squared Error: 0.00

R^2 Score: 1.00



0.1.4 Random Forest Classifier

```
[79]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

[80]: # Define features and target
features = ['Aroma', 'Flavor', 'Aftertaste', 'Acidity', 'Body', 'Balance', \
    'Uniformity', \
    'Clean Cup', 'Sweetness', 'Overall', 'Defects', 'Total Cup Points', \
    'Moisture Percentage', 'Category One Defects', 'Quakers', \
    'Category Two Defects', 'Mean Score']
target = 'Category One Defects' # Replace with your target variable

X = data[features]
y = data[target]
```

```
[81]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[82]: # Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
[83]: # Initialize and train Random Forest Classifier
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train_scaled, y_train)
```

```
[83]: RandomForestClassifier(random_state=42)
```

```
[84]: # Predict on the test set
y_pred = rf_model.predict(X_test_scaled)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(conf_matrix)

# Plot confusion matrix using seaborn
plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=rf_model.classes_, yticklabels=rf_model.classes_)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```

Accuracy: 0.86

Classification Report:

	precision	recall	f1-score	support
0	0.92	1.00	0.96	36
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	3

3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	1
5	0.00	0.00	0.00	1
accuracy			0.86	42
macro avg	0.15	0.17	0.16	42
weighted avg	0.79	0.86	0.82	42

Confusion Matrix:

```
[[36  0  0  0  0  0]
 [ 0  0  0  0  0  0]
 [ 2  1  0  0  0  0]
 [ 1  0  0  0  0  0]
 [ 0  1  0  0  0  0]
 [ 0  1  0  0  0  0]]
```

```
c:\Users\iamim\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.

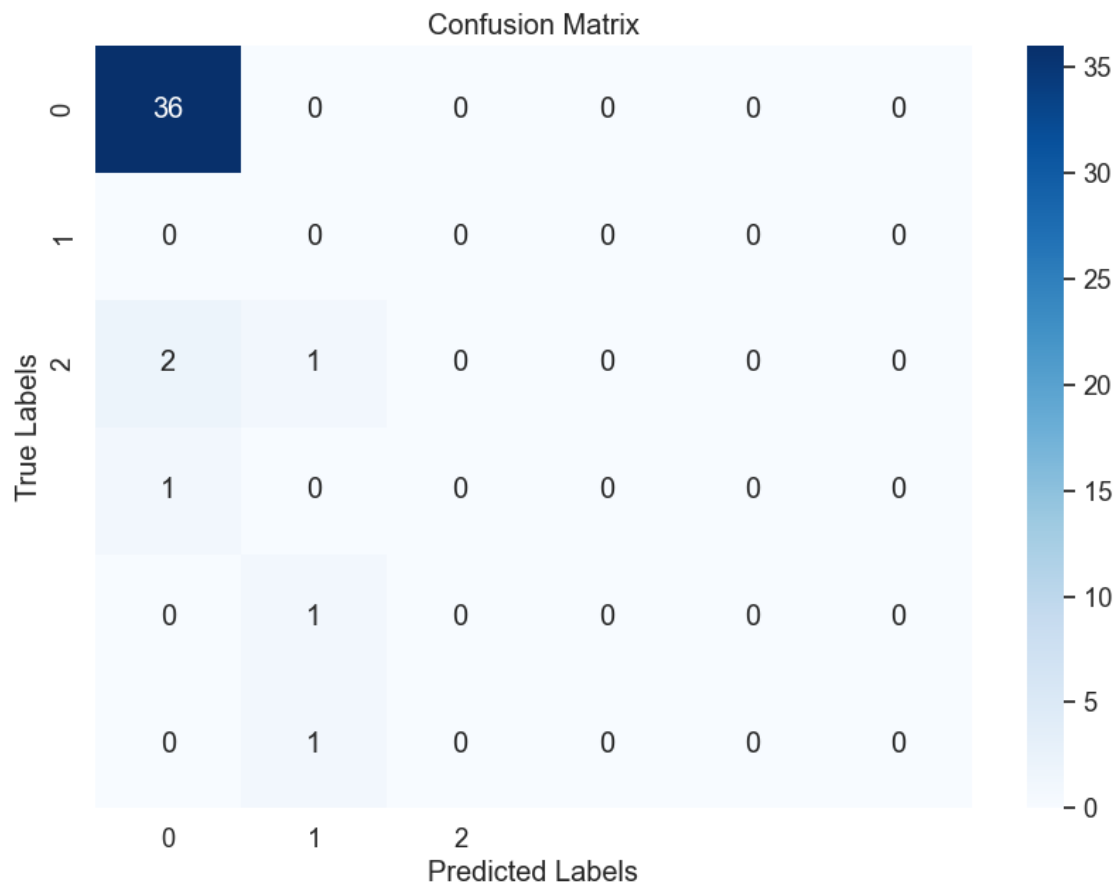
```
c:\Users\iamim\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
c:\Users\iamim\anaconda3\Lib\site-
```

packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning:

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.



0.1.5 Decision Tree Regression

```
[2]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import plot_tree
```

```
[4]: # Load the dataset
```

```
data = pd.read_csv('C:
↳\\Users\\iamim\\OneDrive\\Desktop\\Seventh_Semester\\ML_LAB\\L2\\df_arabica_clean.
↳csv') # Replace 'your_dataset.csv' with your actual file path

# Display the first few rows of the dataset
print(data.head())

# Check for missing values
print(data.isnull().sum())

# Fill missing values with the column mean (if needed)
data.fillna(data.mean(numeric_only=True), inplace=True)
```

	Unnamed: 0	ID	Country of Origin	Farm Name	\
0	0	0	Colombia	Finca El Paraiso	
1	1	1	Taiwan	Royal Bean Geisha Estate	
2	2	2	Laos	OKLAO coffee farms	
3	3	3	Costa Rica	La Cumbre	
4	4	4	Colombia	Finca Santuario	

		Lot Number	Mill	\
0		CQU2022015	Finca El Paraiso	
1	The 2022 Pacific Rim Coffee Summit,T037		Royal Bean Geisha Estate	
2	The 2022 Pacific Rim Coffee Summit,LA01	oklao	coffee processing plant	
3		CQU2022017	La Montana Tarrazu Mill	
4		CQU2023002	Finca Santuario	

	ICO Number	Company	Altitude	Region	...	\
0	NaN	Coffee Quality Union	1700-1930	Piendamo,Cauca	...	
1	NaN	Taiwan Coffee Laboratory	1200	Chiayi	...	
2	NaN	Taiwan Coffee Laboratory	1300	Laos Borofen Plateau	...	
3	NaN	Coffee Quality Union	1900	Los Santos,Tarrazu	...	
4	NaN	Coffee Quality Union	1850-2100	Popayan,Cauca	...	

	Total Cup Points	Moisture Percentage	Category One Defects	Quakers	\
0	89.33	11.8	0	0	
1	87.58	10.5	0	0	
2	87.42	10.4	0	0	
3	87.17	11.8	0	0	
4	87.08	11.6	0	2	

	Color	Category Two Defects	Expiration	\
0	green	3	September 21st, 2023	
1	blue-green	0	November 15th, 2023	
2	yellowish	2	November 15th, 2023	
3	green	0	September 21st, 2023	
4	yellow-green	2	March 5th, 2024	

	Certification Body \
0	Japan Coffee Exchange
1	Taiwan Coffee Laboratory
2	Taiwan Coffee Laboratory
3	Japan Coffee Exchange
4	Japan Coffee Exchange

	Certification Address \
0	413-0002 - 1173-58 Izusan, Ata...
1	QAHWAH CO., LTD 4F, No. 225, Sec. 3, Beixin Rd...
2	QAHWAH CO., LTD 4F, No. 225, Sec. 3, Beixin Rd...
3	413-0002 - 1173-58 Izusan, Ata...
4	413-0002 - 1173-58 Izusan, Ata...

	Certification Contact
0	Koju Matsuzawa - +81(0)9085642901
1	Lin, Jen-An Neil - 886-289116612
2	Lin, Jen-An Neil - 886-289116612
3	Koju Matsuzawa - +81(0)9085642901
4	Koju Matsuzawa - +81(0)9085642901

[5 rows x 41 columns]

Unnamed: 0	0
ID	0
Country of Origin	0
Farm Name	2
Lot Number	1
Mill	3
ICO Number	132
Company	0
Altitude	1
Region	2
Producer	1
Number of Bags	0
Bag Weight	0
In-Country Partner	0
Harvest Year	0
Grading Date	0
Owner	0
Variety	6
Status	0
Processing Method	5
Aroma	0
Flavor	0
Aftertaste	0
Acidity	0
Body	0

```

Balance          0
Uniformity       0
Clean Cup        0
Sweetness        0
Overall          0
Defects          0
Total Cup Points 0
Moisture Percentage 0
Category One Defects 0
Quakers          0
Color           0
Category Two Defects 0
Expiration       0
Certification Body 0
Certification Address 0
Certification Contact 0
dtype: int64

```

```

[5]: # Define features and target variable
features = ['Aroma', 'Flavor', 'Aftertaste'] # Independent variables
target = 'Total Cup Points' # Dependent variable

X = data[features]
y = data[target]

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)

```

```

[6]: # Create the Decision Tree Regressor
dt_regressor = DecisionTreeRegressor(random_state=42)

# Fit the model on the training data
dt_regressor.fit(X_train, y_train)

```

```

[6]: DecisionTreeRegressor(random_state=42)

```

```

[7]: # Predict on the test data
y_pred = dt_regressor.predict(X_test)

```

```

[8]: # Calculate Mean Squared Error (MSE) and R2 Score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

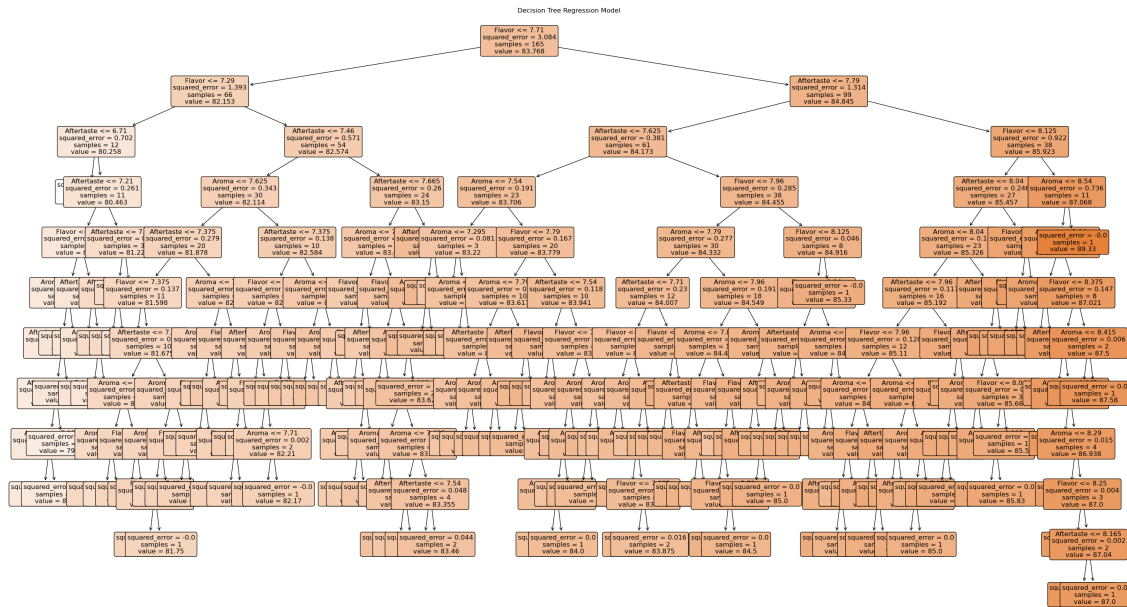
# Print the evaluation metrics
print(f"Mean Squared Error: {mse:.2f}")
print(f"R2 Score: {r2:.2f}")

```

Mean Squared Error: 0.38

R² Score: 0.85

```
[11]: # Plotting the decision tree
plt.figure(figsize=(35, 20))
plot_tree(dt_regressor, feature_names=features, filled=True, rounded=True,
          fontsize=12)
plt.title('Decision Tree Regression Model')
plt.show()
```



```
[12]: # Plot predicted vs actual values
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', edgecolor='k', alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
        linestyle='--', linewidth=2)
plt.xlabel('Actual Total Cup Points')
plt.ylabel('Predicted Total Cup Points')
plt.title('Decision Tree Regression: Predicted vs Actual')
plt.grid(True)
plt.show()
```