

August 14, 2024

## 0.1 ML Assignmnet 4

Name: Imaad Hajwane

SRN: 202101132

Roll No: 23

Program: Computer Engineering

Year: Last year

Div: A

Subject: ML

Q. Write a program to implement bagging and boosting to solve classification problem on datasets.

### 1. Load and Preprocess the Data

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

```
[2]: # Load the dataset
file_path = 'avocado.csv'
df = pd.read_csv(file_path)

# Display the first few rows
print(df.head())
```

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	\
0	0	27-12-2015	1.33	64236.62	1036.74	54454.85	
1	1	20-12-2015	1.35	54876.98	674.28	44638.81	
2	2	13-12-2015	0.93	118220.22	794.70	109149.67	
3	3	06-12-2015	1.08	78992.15	1132.00	71976.41	
4	4	29-11-2015	1.28	51039.60	941.48	43838.39	

	4770	Total Bags	Small Bags	Large Bags	XLarge Bags	type	\
0	48.16	8696.87	8603.62	93.25	0.0	conventional	
1	58.33	9505.56	9408.07	97.49	0.0	conventional	
2	130.50	8145.35	8042.21	103.14	0.0	conventional	
3	72.58	5811.16	5677.40	133.76	0.0	conventional	

```
4    75.78    6183.95    5986.26    197.69    0.0    conventional
```

```
   year region
0  2015  Albany
1  2015  Albany
2  2015  Albany
3  2015  Albany
4  2015  Albany
```

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      18249 non-null  int64
1   Date            18249 non-null  object
2   AveragePrice    18249 non-null  float64
3   Total Volume    18249 non-null  float64
4   4046            18249 non-null  float64
5   4225            18249 non-null  float64
6   4770            18249 non-null  float64
7   Total Bags      18249 non-null  float64
8   Small Bags      18249 non-null  float64
9   Large Bags      18249 non-null  float64
10  XLarge Bags     18249 non-null  float64
11  type            18249 non-null  object
12  year            18249 non-null  int64
13  region          18249 non-null  object
dtypes: float64(9), int64(2), object(3)
memory usage: 1.9+ MB
```

```
[4]: df.isnull().sum()
```

```
[4]: Unnamed: 0      0
     Date          0
     AveragePrice  0
     Total Volume  0
     4046          0
     4225          0
     4770          0
     Total Bags    0
     Small Bags    0
     Large Bags    0
     XLarge Bags   0
     type          0
```

```
year            0
region          0
dtype: int64
```

```
[5]: # Preprocess the data
# Convert 'Date' to datetime format
df['Date'] = pd.to_datetime(df['Date'])

# Encode categorical variables
label_encoders = {}
for column in ['type', 'region']:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le

# Drop any rows with missing values (if any)
df.dropna(inplace=True)

# Separate features and target variable
X = df.drop(['AveragePrice', 'Date'], axis=1)
y = df['AveragePrice']

# Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

C:\Users\iamim\AppData\Local\Temp\ipykernel\_24820\810139836.py:3: UserWarning:  
Parsing dates in %d-%m-%Y format when dayfirst=False (the default) was  
specified. Pass `dayfirst=True` or specify a format to silence this warning.

```
df['Date'] = pd.to_datetime(df['Date'])
```

## 2. Train-Test Split

## 3. Implement the KNN Classifier

```
[10]: # Convert 'AveragePrice' to a binary classification target (1 if above median, 0 if below)
median_price = df['AveragePrice'].median()
df['PriceCategory'] = (df['AveragePrice'] > median_price).astype(int)

# Separate features and new target variable
X_classification = df.drop(['AveragePrice', 'Date', 'PriceCategory'], axis=1)
y_classification = df['PriceCategory']

# Scale the features
X_scaled_classification = scaler.fit_transform(X_classification)
```

```
[11]: # Split the data into training and testing sets
X_train_clf, X_test_clf, y_train_clf, y_test_clf = \
    train_test_split(X_scaled_classification, y_classification, test_size=0.2,
                    random_state=42)
```

```
[12]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report, \
    accuracy_score, roc_curve, auc

# Initialize the KNN Classifier
knn_clf = KNeighborsClassifier(n_neighbors=5)

# Train the model
knn_clf.fit(X_train_clf, y_train_clf)

# Make predictions
y_pred_clf = knn_clf.predict(X_test_clf)

# Evaluate the model
accuracy = accuracy_score(y_test_clf, y_pred_clf)
conf_matrix = confusion_matrix(y_test_clf, y_pred_clf)
class_report = classification_report(y_test_clf, y_pred_clf)

# Print the evaluation metrics
print(f'Accuracy: {accuracy}')
print('Confusion Matrix:')
print(conf_matrix)
print('Classification Report:')
print(class_report)
```

Accuracy: 0.8780821917808219

Confusion Matrix:

```
[[1526  267]
```

```
 [ 178 1679]]
```

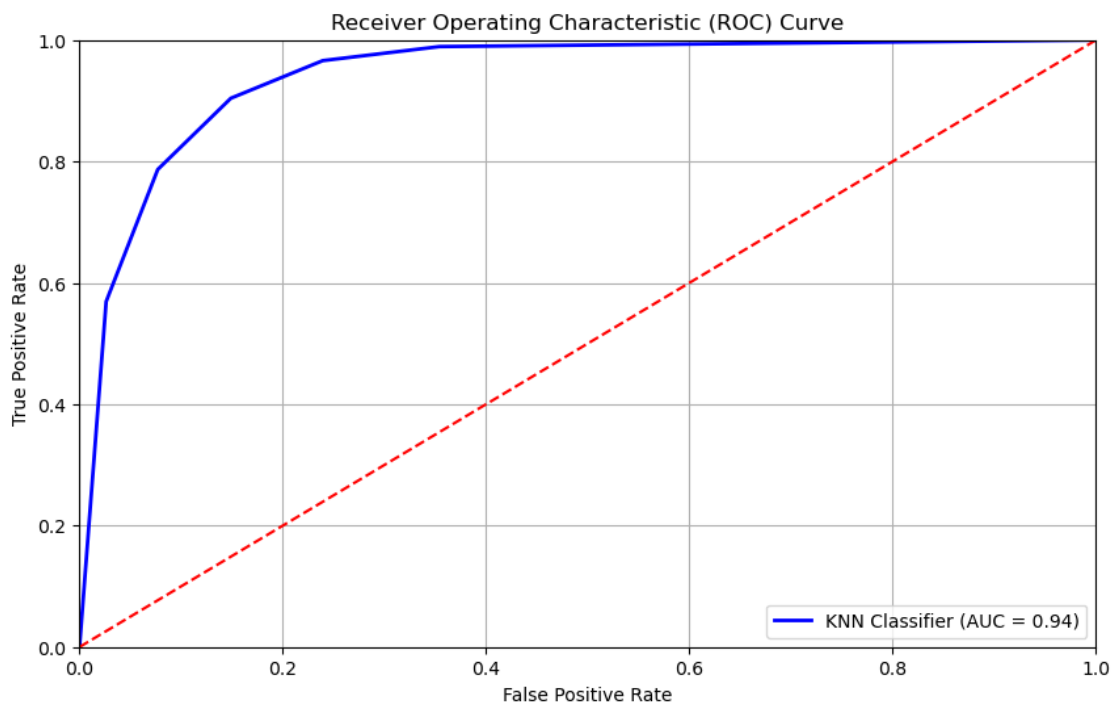
Classification Report:

	precision	recall	f1-score	support
0	0.90	0.85	0.87	1793
1	0.86	0.90	0.88	1857
accuracy			0.88	3650
macro avg	0.88	0.88	0.88	3650
weighted avg	0.88	0.88	0.88	3650

```
[16]: # Calculate the ROC curve and AUC
y_pred_prob = knn_clf.predict_proba(X_test_clf)[: , 1]
```

```
fpr, tpr, thresholds = roc_curve(y_test_clf, y_pred_prob)
roc_auc = auc(fpr, tpr)


# Plot the ROC curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'KNN Classifier (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.grid(True)
plt.savefig('knn_roc_curve.png')
plt.show()
```



```
[18]: import seaborn as sns
import matplotlib.pyplot as plt

# Confusion Matrix Plot
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
```


```
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.savefig('knn_confusion_matrix.png', dpi=300)
plt.show()
```



```
[20]: from sklearn.metrics import precision_recall_curve

# Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test_clf, y_pred_prob)

plt.figure(figsize=(10, 6))
plt.plot(recall, precision, color='purple', lw=2)
plt.title('Precision-Recall Curve')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.grid(True)
plt.savefig('knn_precision_recall_curve.png', dpi=300)
plt.show()
```

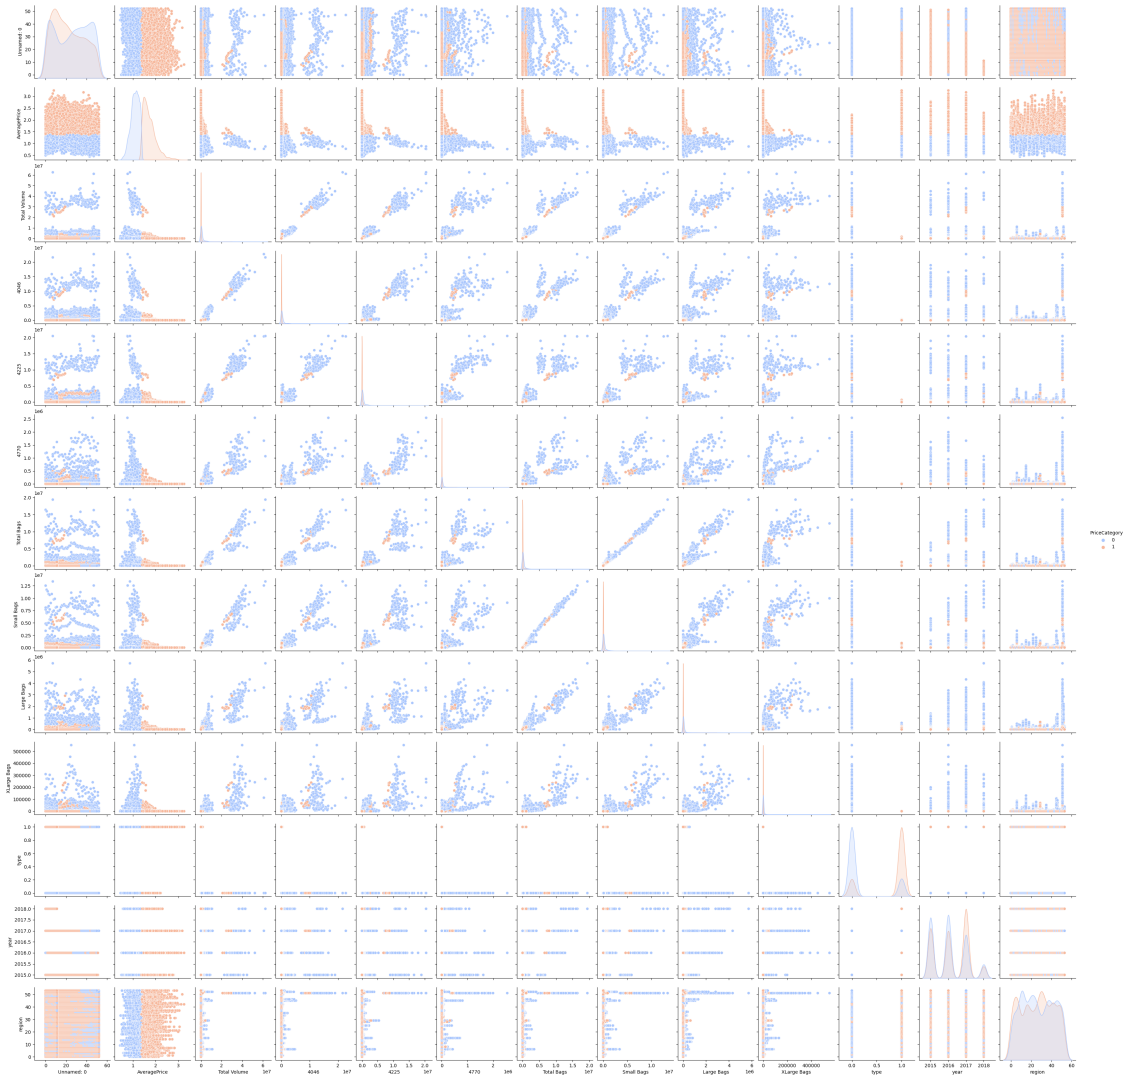


### Pairplot of Features

```
[21]: sns.pairplot(df, hue='PriceCategory', palette='coolwarm', diag_kind='kde')  
plt.savefig('knn_pairplot.png', dpi=300)  
plt.show()
```

c:\Users\iamim\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning:

The figure layout has changed to tight



```
[22]: import plotly.express as px

# Confusion Matrix in Plotly
fig_conf_matrix = px.imshow(conf_matrix, text_auto=True, labels={'x': 'Predicted', 'y': 'True'}, title='Confusion Matrix')
fig_conf_matrix.write_html('knn_confusion_matrix.html')

# ROC Curve in Plotly
fig_roc = px.area(x=fpr, y=tpr, title=f'ROC Curve (AUC = {roc_auc:.2f})',
                  labels={'x': 'False Positive Rate', 'y': 'True Positive Rate'},
                  hover_data={'False Positive Rate': fpr, 'True Positive Rate': tpr})
fig_roc.add_shape(type='line', line=dict(dash='dash'), x0=0, x1=1, y0=0, y1=1)
```




```
fig_roc.write_html('knn_roc_curve.html')

# Precision-Recall Curve in Plotly
fig_pr = px.line(x=recall, y=precision, title='Precision-Recall Curve',
                 labels={'x': 'Recall', 'y': 'Precision'})
fig_pr.write_html('knn_precision_recall_curve.html')
```

### Distribution Plot of AveragePrice


```
[23]: plt.figure(figsize=(10, 6))
sns.histplot(df['AveragePrice'], kde=True, color='blue')
plt.title('Distribution of Average Price')
plt.xlabel('Average Price')
plt.ylabel('Frequency')
plt.grid(True)
plt.savefig('average_price_distribution.png', dpi=300)
plt.show()
```



### Correlation Heatmap

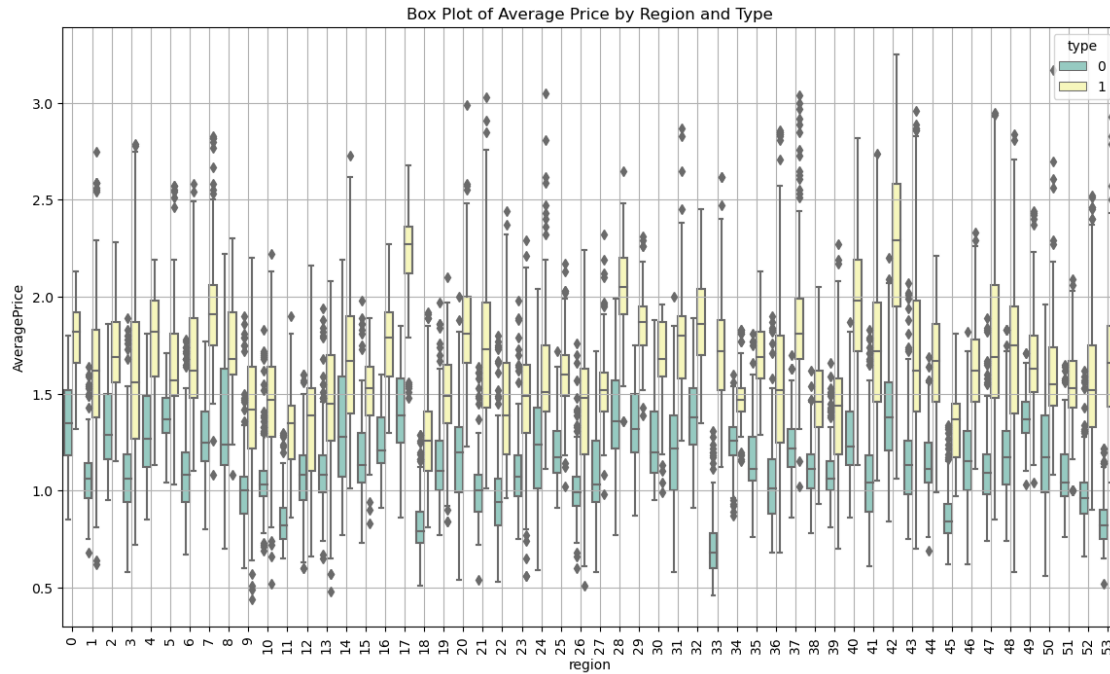
```
[24]: plt.figure(figsize=(12, 8))
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
            linewidths=0.5)
```

```
plt.title('Correlation Heatmap')
plt.savefig('correlation_heatmap.png', dpi=300)
plt.show()
```




### Box Plot of AveragePrice by Region and Type

```
[25]: plt.figure(figsize=(14, 8))
sns.boxplot(x='region', y='AveragePrice', hue='type', data=df, palette='Set3')
plt.title('Box Plot of Average Price by Region and Type')
plt.xticks(rotation=90)
plt.grid(True)
plt.savefig('boxplot_average_price_region_type.png', dpi=300)
plt.show()
```



### Time Series Plot of AveragePrice

```
[26]: plt.figure(figsize=(14, 8))
sns.lineplot(x='Date', y='AveragePrice', hue='type', data=df,
             palette='coolwarm')
plt.title('Average Price Over Time')
plt.xlabel('Date')
plt.ylabel('Average Price')
plt.grid(True)
plt.savefig('average_price_over_time.png', dpi=300)
plt.show()
```



```
[27]: # Distribution Plot using Plotly
fig_dist = px.histogram(df, x='AveragePrice', nbins=50, marginal='box',
    ↪title='Distribution of Average Price')
fig_dist.write_html('average_price_distribution.html')

# Correlation Heatmap using Plotly
fig_corr = px.imshow(correlation_matrix, text_auto=True, aspect='auto',
    ↪title='Correlation Heatmap')
fig_corr.write_html('correlation_heatmap.html')

# Box Plot using Plotly
fig_box = px.box(df, x='region', y='AveragePrice', color='type', title='Box
    ↪Plot of Average Price by Region and Type')
fig_box.update_layout(xaxis_tickangle=-90)
fig_box.write_html('boxplot_average_price_region_type.html')

# Time Series Plot using Plotly
fig_time = px.line(df, x='Date', y='AveragePrice', color='type', title='Average
    ↪Price Over Time')
fig_time.write_html('average_price_over_time.html')
```

---

Implementation of Ada boost, Gradient boosting & Random Forest, Bagging classifiers

```
[28]: from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier, \
      ↪ RandomForestClassifier, BaggingClassifier
      from sklearn.metrics import accuracy_score, classification_report
      import matplotlib.pyplot as plt
      import seaborn as sns

      # Initialize a dictionary to store the accuracy of each model
      model_accuracies = {}
```

AdaBoost Classifier

```
[29]: # AdaBoost Classifier
      ada_clf = AdaBoostClassifier(n_estimators=50, random_state=42)
      ada_clf.fit(X_train_clf, y_train_clf)
      y_pred_ada = ada_clf.predict(X_test_clf)

      # Accuracy and classification report
      ada_accuracy = accuracy_score(y_test_clf, y_pred_ada)
      model_accuracies['AdaBoost'] = ada_accuracy
      print(f"AdaBoost Accuracy: {ada_accuracy:.2f}")
      print(classification_report(y_test_clf, y_pred_ada))
```

AdaBoost Accuracy: 0.83

	precision	recall	f1-score	support
0	0.81	0.84	0.83	1793
1	0.84	0.81	0.83	1857
accuracy			0.83	3650
macro avg	0.83	0.83	0.83	3650
weighted avg	0.83	0.83	0.83	3650

Gradient Boosting Classifier

```
[30]: # Gradient Boosting Classifier
      gb_clf = GradientBoostingClassifier(n_estimators=100, random_state=42)
      gb_clf.fit(X_train_clf, y_train_clf)
      y_pred_gb = gb_clf.predict(X_test_clf)

      # Accuracy and classification report
      gb_accuracy = accuracy_score(y_test_clf, y_pred_gb)
      model_accuracies['Gradient Boosting'] = gb_accuracy
      print(f"Gradient Boosting Accuracy: {gb_accuracy:.2f}")
      print(classification_report(y_test_clf, y_pred_gb))
```

Gradient Boosting Accuracy: 0.87

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.87	0.87	0.87	1793
1	0.87	0.87	0.87	1857
accuracy				0.87 3650
macro avg	0.87	0.87	0.87	3650
weighted avg	0.87	0.87	0.87	3650

#### Random Forest Classifier

```
[31]: # Random Forest Classifier
rf_clf = RandomForestClassifier(n_estimators=100, random_state=42)
rf_clf.fit(X_train_clf, y_train_clf)
y_pred_rf = rf_clf.predict(X_test_clf)

# Accuracy and classification report
rf_accuracy = accuracy_score(y_test_clf, y_pred_rf)
model_accuracies['Random Forest'] = rf_accuracy
print(f"Random Forest Accuracy: {rf_accuracy:.2f}")
print(classification_report(y_test_clf, y_pred_rf))
```

Random Forest Accuracy: 0.93

	precision	recall	f1-score	support
0	0.93	0.92	0.92	1793
1	0.92	0.94	0.93	1857
accuracy				0.93 3650
macro avg	0.93	0.93	0.93	3650
weighted avg	0.93	0.93	0.93	3650

#### Bagging Classifier

```
[32]: # Bagging Classifier
bagging_clf = BaggingClassifier(base_estimator=KNeighborsClassifier(),
    ↪n_estimators=50, random_state=42)
bagging_clf.fit(X_train_clf, y_train_clf)
y_pred_bagging = bagging_clf.predict(X_test_clf)

# Accuracy and classification report
bagging_accuracy = accuracy_score(y_test_clf, y_pred_bagging)
model_accuracies['Bagging'] = bagging_accuracy
print(f"Bagging Classifier Accuracy: {bagging_accuracy:.2f}")
print(classification_report(y_test_clf, y_pred_bagging))
```

c:\Users\iamim\anaconda3\Lib\site-packages\sklearn\ensemble\\_base.py:156:  
FutureWarning:

`base\_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.

Bagging Classifier Accuracy: 0.88

	precision	recall	f1-score	support
0	0.90	0.85	0.88	1793
1	0.86	0.91	0.89	1857
accuracy			0.88	3650
macro avg	0.88	0.88	0.88	3650
weighted avg	0.88	0.88	0.88	3650

### Plot the Accuracies Together

```
[42]: import matplotlib.pyplot as plt
import seaborn as sns

# Example model accuracies (replace with your actual data)
model_accuracies = {
    'KNN': 0.88,
    'AdaBoost': 0.83,
    'Gradient Boosting': 0.87,
    'Random Forest': 0.93,
    'Bagging Classifier': 0.88
}

# Professional Plot for Model Comparison
plt.figure(figsize=(14, 8))

# Bar plot for model accuracies
sns.barplot(x=list(model_accuracies.keys()), y=list(model_accuracies.values()),
            palette='dark:#5A9_r', alpha=0.85, edgecolor='black')

# Line plot for model accuracies
sns.lineplot(x=list(model_accuracies.keys()), y=list(model_accuracies.values()),
             marker='o', color='black', linewidth=2.5, markersize=10)

# Adding text annotations for each bar
for i, v in enumerate(model_accuracies.values()):
    plt.text(i, v + 0.005, f"{v:.2f}", ha='center', fontsize=12,
            fontweight='bold', color='black')

plt.title('Model Comparison - Accuracy', fontsize=18, fontweight='bold',
        color='#333333')
plt.ylabel('Accuracy', fontsize=14, fontweight='bold', color='#333333')
```

```

plt.xlabel('Model', fontsize=14, fontweight='bold', color='#333333')

# Adjusting y-axis limits to ensure all values are visible
plt.ylim(0.80, 1.0) # Increased lower limit to 0.80

plt.xticks(fontsize=12, fontweight='bold', color='#333333')
plt.yticks(fontsize=12, fontweight='bold', color='#333333')


# Adding subtle grid lines
plt.grid(True, which='major', linestyle='--', linewidth=0.6, color='grey',
        alpha=0.7)

# Remove top and right borders
sns.despine()

plt.tight_layout()

# Save the plot with higher DPI for better quality
plt.savefig('professional_model_comparison_accuracy.png', dpi=300,
        bbox_inches='tight')
plt.show()

```



## Displaying Classification Reports

```

[44]: import pandas as pd
      from sklearn.metrics import classification_report, accuracy_score

```



```

# Example predictions and true labels (replace with your actual data)
# y_test_clf, y_pred_clf, y_pred_ada, y_pred_gb, y_pred_rf, y_pred_bagging
↳ should be defined

# Store the classification reports and accuracies in dictionaries
classification_reports = {
    'KNN': classification_report(y_test_clf, y_pred_clf, output_dict=True),
    'AdaBoost': classification_report(y_test_clf, y_pred_ada, output_dict=True),
    'Gradient Boosting': classification_report(y_test_clf, y_pred_gb,
↳ output_dict=True),
    'Random Forest': classification_report(y_test_clf, y_pred_rf,
↳ output_dict=True),
    'Bagging': classification_report(y_test_clf, y_pred_bagging,
↳ output_dict=True)
}

accuracies = {
    'KNN': accuracy_score(y_test_clf, y_pred_clf),
    'AdaBoost': accuracy_score(y_test_clf, y_pred_ada),
    'Gradient Boosting': accuracy_score(y_test_clf, y_pred_gb),
    'Random Forest': accuracy_score(y_test_clf, y_pred_rf),
    'Bagging': accuracy_score(y_test_clf, y_pred_bagging)
}

# Convert the reports into DataFrames
reports_df = {model: pd.DataFrame(report).transpose() for model, report in
↳ classification_reports.items()}

# Add accuracy as a new row to each DataFrame
for model, df in reports_df.items():
    df.loc['accuracy', :] = pd.Series(accuracies[model], index=df.columns)

# Add accuracy to the classification report for clarity
for model, df in reports_df.items():
    df.loc['accuracy', :] = pd.Series(accuracies[model], index=df.columns)

# Create a DataFrame for accuracies
accuracy_df = pd.DataFrame(list(accuracies.items()), columns=['Model',
↳ 'Accuracy']).set_index('Model')

# Combine the DataFrames into a single DataFrame for better visualization
combined_report_df = pd.concat(reports_df, axis=1)

# Append the accuracies DataFrame to the combined report DataFrame
combined_report_df = pd.concat([combined_report_df, accuracy_df.T], axis=0)

```

```
# Display the combined classification report
with pd.option_context('display.max_columns', None): # to display all columns
    print(combined_report_df)
```

	(KNN, precision)	(KNN, recall)	(KNN, f1-score)	\
0	0.895540	0.851088	0.872748	
1	0.862795	0.904146	0.882987	
accuracy	0.878082	0.878082	0.878082	
macro avg	0.879168	0.877617	0.877868	
weighted avg	0.878881	0.878082	0.877957	
Accuracy	NaN	NaN	NaN	

	(KNN, support)	(AdaBoost, precision)	(AdaBoost, recall)	\
0	1793.000000	0.813678	0.842722	
1	1857.000000	0.842722	0.813678	
accuracy	0.878082	0.827945	0.827945	
macro avg	3650.000000	0.828200	0.828200	
weighted avg	3650.000000	0.828454	0.827945	
Accuracy	NaN	NaN	NaN	

	(AdaBoost, f1-score)	(AdaBoost, support)	\
0	0.827945	1793.000000	
1	0.827945	1857.000000	
accuracy	0.827945	0.827945	
macro avg	0.827945	3650.000000	
weighted avg	0.827945	3650.000000	
Accuracy	NaN	NaN	

	(Gradient Boosting, precision)	(Gradient Boosting, recall)	\
0	0.868083	0.866146	
1	0.871037	0.872913	
accuracy	0.869589	0.869589	
macro avg	0.869560	0.869530	
weighted avg	0.869586	0.869589	
Accuracy	NaN	NaN	

	(Gradient Boosting, f1-score)	(Gradient Boosting, support)	\
0	0.867113	1793.000000	
1	0.871974	1857.000000	
accuracy	0.869589	0.869589	
macro avg	0.869544	3650.000000	
weighted avg	0.869586	3650.000000	
Accuracy	NaN	NaN	

	(Random Forest, precision)	(Random Forest, recall)	\
0	0.932955	0.915784	

1	0.920106	0.936457
accuracy	0.926301	0.926301
macro avg	0.926530	0.926120
weighted avg	0.926418	0.926301
Accuracy	NaN	NaN

	(Random Forest, f1-score)	(Random Forest, support) \
0	0.924289	1793.000000
1	0.928209	1857.000000
accuracy	0.926301	0.926301
macro avg	0.926249	3650.000000
weighted avg	0.926284	3650.000000
Accuracy	NaN	NaN

	(Bagging, precision)	(Bagging, recall)	(Bagging, f1-score) \
0	0.900885	0.851645	0.875573
1	0.863939	0.909532	0.886149
accuracy	0.881096	0.881096	0.881096
macro avg	0.882412	0.880588	0.880861
weighted avg	0.882088	0.881096	0.880954
Accuracy	NaN	NaN	NaN

	(Bagging, support)	KNN	AdaBoost	Gradient Boosting \
0	1793.000000	NaN	NaN	NaN
1	1857.000000	NaN	NaN	NaN
accuracy	0.881096	NaN	NaN	NaN
macro avg	3650.000000	NaN	NaN	NaN
weighted avg	3650.000000	NaN	NaN	NaN
Accuracy	NaN	0.878082	0.827945	0.869589

	Random Forest	Bagging
0	NaN	NaN
1	NaN	NaN
accuracy	NaN	NaN
macro avg	NaN	NaN
weighted avg	NaN	NaN
Accuracy	0.926301	0.881096